

# Spotify Hit Prediction

Let's try to predict the song will be hit or miss.

This spotify dataset has songs from 1960s-2010s.



## All About Data

```
In [1]: import numpy as np
import seaborn as sns
from sklearn.ensemble import RandomForestClassifier, VotingClassifier, StackingClassifier,
import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
from sklearn.linear_model import LogisticRegression
from sklearn.model_selection import cross_val_score, RepeatedStratifiedKFold
from sklearn.svm import LinearSVC, SVC
from sklearn.ensemble import RandomForestClassifier, GradientBoostingClassifier
import warnings
from sklearn.metrics import confusion_matrix
from sklearn import metrics
warnings.filterwarnings(action='ignore')
import matplotlib.pyplot as plt
import seaborn as sns
```

```
In [2]: datas = [pd.read_csv("/content/drive/MyDrive/Colab Notebooks/AppliedML/archive/dataset-of-
```

```
In [3]: for i, decade in enumerate([1960, 1970, 1980, 1990, 2000, 2010]):
    datas[i]['decade'] = pd.Series(decade, index=datas[i].index)

data = pd.concat(datas, axis=0).sample(frac=1.0, random_state=1).reset_index(drop=True)
```

```
In [4]: data.head()
```

```
Out[4]:
```

	track	artist	uri	danceability	energy	key	loudness	mode	speechi
0	Attaining - Take 1 / Alternate Version	John Coltrane	spotify:track:3EwLV5hZqLKx5e0Lp1QcB7	0.342	0.462	4	-12.931	0	0.0

	track	artist	uri	danceability	energy	key	loudness	mode	speechiness
1	So Fly	NB Ridaz Featuring Gemini	spotify:track:2Bjli07kN0yKSur0Fwrnss	0.861	0.519	2	-6.404	1	0.1
2	Because I Got It Like That	Jungle Brothers	spotify:track:5unLExF3iiG3YkU11u6wFO	0.900	0.916	1	-7.481	0	0.1
3	Babylon a Fall - Remastered	Yabby You	spotify:track:6xfe0G2HwRDQaChxkzvNKw	0.714	0.301	2	-14.800	1	0.1
4	Fins	Jimmy Buffett	spotify:track:4h0gZ422QxBRdTV14u0P8y	0.661	0.645	4	-13.520	1	0.0

In [5]: `data.shape`

Out[5]: (41106, 20)

Data has 41106 rows and 20 columns.

In [6]: `data.info()`

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 41106 entries, 0 to 41105
Data columns (total 20 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   track                                41106 non-null  object
1   artist                              41106 non-null  object
2   uri                                  41106 non-null  object
3   danceability                        41106 non-null  float64
4   energy                              41106 non-null  float64
5   key                                  41106 non-null  int64
6   loudness                            41106 non-null  float64
7   mode                                41106 non-null  int64
8   speechiness                         41106 non-null  float64
9   acousticness                       41106 non-null  float64
10  instrumentalness                    41106 non-null  float64
11  liveness                            41106 non-null  float64
12  valence                             41106 non-null  float64
13  tempo                               41106 non-null  float64
14  duration_ms                         41106 non-null  int64
15  time_signature                      41106 non-null  int64
16  chorus_hit                          41106 non-null  float64
17  sections                            41106 non-null  int64
18  target                              41106 non-null  int64
19  decade                             41106 non-null  int64
dtypes: float64(10), int64(7), object(3)
memory usage: 6.3+ MB
```

In [7]: `data.columns`

Out[7]: Index(['track', 'artist', 'uri', 'danceability', 'energy', 'key', 'loudness', 'mode', 'speechiness', 'acousticness', 'instrumentalness', 'liveness', 'valence', 'tempo', 'duration\_ms', 'time\_signature', 'chorus\_hit', 'sections', 'target', 'decade'], dtype='object')

In [8]:

```
data.nunique(axis=0)
```

```
Out[8]: track          35860
artist        11904
uri           40560
danceability   1048
energy         1787
key            12
loudness      16160
mode           2
speechiness    1346
acousticness   4194
instrumentalness 5122
liveness       1674
valence        1609
tempo         32152
duration_ms    21517
time_signature 5
chorus_hit     39950
sections       84
target         2
decade         6
dtype: int64
```

```
In [9]: data.describe().apply(lambda s: s.apply(lambda x: format(x, 'f')))
```

```
Out[9]:
```

	danceability	energy	key	loudness	mode	speechiness	acousticness	instrumen
<b>count</b>	41106.000000	41106.000000	41106.000000	41106.000000	41106.000000	41106.000000	41106.000000	41106.
<b>mean</b>	0.539695	0.579545	5.213594	-10.221525	0.693354	0.072960	0.364197	0.
<b>std</b>	0.177821	0.252628	3.534977	5.311626	0.461107	0.086112	0.338913	0.
<b>min</b>	0.000000	0.000251	0.000000	-49.253000	0.000000	0.000000	0.000000	0.
<b>25%</b>	0.420000	0.396000	2.000000	-12.816000	0.000000	0.033700	0.039400	0.
<b>50%</b>	0.552000	0.601000	5.000000	-9.257000	1.000000	0.043400	0.258000	0.
<b>75%</b>	0.669000	0.787000	8.000000	-6.374250	1.000000	0.069800	0.676000	0.
<b>max</b>	0.988000	1.000000	11.000000	3.744000	1.000000	0.960000	0.996000	1.

```
In [10]: total = data.isnull().sum().sort_values(ascending=False)
percent = (data.isnull().sum()/data.isnull().count()).sort_values(ascending=False)
missing_data = pd.concat([total,percent],axis=1,keys=["total","percent"])
missing_data.head()
```

```
Out[10]:
```

	total	percent
<b>track</b>	0	0.0
<b>artist</b>	0	0.0
<b>target</b>	0	0.0
<b>sections</b>	0	0.0
<b>chorus_hit</b>	0	0.0

There are no missing values in the Data.

Let's check how many categorical and numerical values are present in the data.

```
In [11]: len(data._get_numeric_data().columns)
```

```
Out[11]: 17
```

There are 17 numeric columns and 3 categorical columns.

```
In [12]: categorical_cols=data.columns[data.dtypes =='object']
print(categorical_cols)
```

```
Index(['track', 'artist', 'uri'], dtype='object')
```

## Data Preprocessing

We are performing following steps in preprocessing:

- Removing categorical variables
- Standard scaling of the data
- Splitted the data in 70%, 15%, 15% as train, validation and test dataset

```
In [13]: def preprocessing(data_df):
        data_prev = data_df.copy()

        """ Let's drop the categorical columns for our analysis
        """
        data_df = data_df.drop(['track', 'artist', 'uri'], axis=1)

        y = data_df['target']
        X = data_df.drop('target', axis=1)
        print(X.shape,y.shape)

        """ Splitting of data
        """
        X_inter, X_test, y_inter, y_test = train_test_split(X, y, train_size=0.8, test_size=0.2)

        X_train, X_val, y_train, y_val = train_test_split(X_inter, y_inter, train_size=0.75, test_size=0.25)

        """ Standard Scaling of data
        """
        scaler = StandardScaler()
        """ Only passing training set to avoid data leakage
        """
        scaler.fit(X_train)
        X_train = pd.DataFrame(scaler.transform(X_train), index=X_train.index, columns=X_train.columns)
        X_val = pd.DataFrame(scaler.transform(X_val), index=X_val.index, columns=X_val.columns)
        X_test = pd.DataFrame(scaler.transform(X_test), index=X_test.index, columns=X_test.columns)

        return X_train, X_test, X_val, y_train, y_test, y_val
```

```
In [14]: X_train, X_test, X_val, y_train, y_test, y_val = preprocessing(data)

(41106, 16) (41106,)
```

```
In [15]: print(X_train.shape)
print(X_test.shape)
print(X_val.shape)

(24663, 16)
```

```
(8222, 16)
(8221, 16)
```

In [16]:

```
print(y_train.shape)
print(y_test.shape)
print(y_val.shape)
```

```
(24663,)
(8222,)
(8221,)
```

# Model Training

## Logistic Regression (softmax regression)

In [17]:

```
def Logistic_regression(solver,max_iter=100,C=1,penalty=None):
    model = LogisticRegression(solver=solver,C=C, max_iter=max_iter,penalty=None)
    model.fit(X_train, y_train)
    n_scores_val = model.score(X_val, y_val)
    n_scores_train = model.score(X_train, y_train)
    print('Mean training Accuracy:',n_scores_train)
    print('Mean validation Accuracy:',n_scores_val)
    return n_scores_train, n_scores_val
```

In [18]:

```
solver_list=["lbfgs", "newton-cg", "sag", "saga"]
for e in solver_list:
    print("Hyper parameter - Solver: ", e, "\n")
    Logistic_regression(e)
    print("----- \n")
```

Hyper parameter - Solver: lbfgs

Mean training Accuracy: 0.7421643757855898  
Mean validation Accuracy: 0.7404208733730689  
-----

Hyper parameter - Solver: newton-cg

Mean training Accuracy: 0.7421643757855898  
Mean validation Accuracy: 0.7404208733730689  
-----

Hyper parameter - Solver: sag

Mean training Accuracy: 0.7421643757855898  
Mean validation Accuracy: 0.7404208733730689  
-----

Hyper parameter - Solver: saga

Mean training Accuracy: 0.7421643757855898  
Mean validation Accuracy: 0.7404208733730689  
-----

In [19]:

```
# Using hyper parameter - solver ("lbfgs", "newton-cg", "sag", "saga") and max_iteration=
solver_list=["lbfgs", "newton-cg", "sag", "saga"]
for e in solver_list:
    print("Hyper parameter - Solver: ", e, "\n")
```

```
Logistic_regression(solver=e,max_iter=1000,C=0.2)
print("----- \n")
```

Hyper parameter - Solver: lbfgs

Mean training Accuracy: 0.7421643757855898  
Mean validation Accuracy: 0.7404208733730689

Hyper parameter - Solver: newton-cg

Mean training Accuracy: 0.7421643757855898  
Mean validation Accuracy: 0.7404208733730689

Hyper parameter - Solver: sag

Mean training Accuracy: 0.7421643757855898  
Mean validation Accuracy: 0.7404208733730689

Hyper parameter - Solver: saga

Mean training Accuracy: 0.7421643757855898  
Mean validation Accuracy: 0.7404208733730689

In [20]:

```
# Using hyper parameter - solver ("lbfgs", "sag", "saga") and max_iteration=100
solver_list=["lbfgs", "sag", "saga"]
for e in solver_list:
    print("Hyper parameter - Solver: ", e)
    Logistic_regression(solver=e,C=10)
    print("\n")
    print("----- \n")
```

Hyper parameter - Solver: lbfgs

Mean training Accuracy: 0.7421643757855898  
Mean validation Accuracy: 0.7404208733730689

Hyper parameter - Solver: sag

Mean training Accuracy: 0.7421643757855898  
Mean validation Accuracy: 0.7404208733730689

Hyper parameter - Solver: saga

Mean training Accuracy: 0.7421643757855898  
Mean validation Accuracy: 0.7404208733730689

In [21]:

```
# Using hyper parameter - solver ("lbfgs", "newton-cg", "sag", "saga") and C=0.8
solver_list=["lbfgs", "newton-cg", "sag", "saga"]
for e in solver_list:
    print("Hyper parameter - Solver: ", e, "\n")
    Logistic_regression(solver=e,C=0.8)
    print("----- \n")
```

```
Hyper parameter - Solver:  lbfgs

Mean training Accuracy: 0.7421643757855898
Mean validation Accuracy: 0.7404208733730689
-----

Hyper parameter - Solver:  newton-cg

Mean training Accuracy: 0.7421643757855898
Mean validation Accuracy: 0.7404208733730689
-----

Hyper parameter - Solver:  sag

Mean training Accuracy: 0.7421643757855898
Mean validation Accuracy: 0.7404208733730689
-----

Hyper parameter - Solver:  saga

Mean training Accuracy: 0.7421643757855898
Mean validation Accuracy: 0.7404208733730689
-----
```

In [43]:

```
penalty=['l1', 'l2', 'elasticnet']
for e in penalty:
    print("Hyper parameter - Penalty: ", e, "\n")
    Logistic_regression(solver="lbfgs",penalty=e,C=0.0001)
    print("----- \n")
```

```
Hyper parameter - Penalty:  l1

Mean training Accuracy: 0.7421643757855898
Mean validation Accuracy: 0.7404208733730689
-----

Hyper parameter - Penalty:  l2

Mean training Accuracy: 0.7421643757855898
Mean validation Accuracy: 0.7404208733730689
-----

Hyper parameter - Penalty:  elasticnet

Mean training Accuracy: 0.7421643757855898
Mean validation Accuracy: 0.7404208733730689
-----
```

### Observation:

Looking at the results seems that the choice of solver hyperparameter does not have a significant impact on the performance of logistic regression for the given dataset and problem.

All four solvers (lbfgs, newton-cg, sag, and saga) gave similar mean training and validation accuracies, with no clear indication that any one solver is better than the others. Therefore, it may be best to choose the solver based on other considerations, such as the computational efficiency or suitability for the problem at hand.

Regardless of the choice of solver, regularization penalty, or maximum iterations, the mean training and validation accuracies remained consistent at around 0.74. Similarly, changing the value of C or using different penalties (l1, l2, elasticnet) did not result in any noticeable improvement in the model's performance.

One possible explanation for the lack of significant improvement is that the dataset may be relatively simple or linearly separable, and therefore logistic regression with default hyperparameters is already able to capture the underlying patterns well.

Another possible reason is model has already converged to the optimal solution and further hyperparameter tuning won't result in any significant improvement. In such cases, changing the hyperparameters may not cause significant changes in the model's accuracy.

## Support Vector machines with hyperparameter tuning

In [23]:

```
C_list = [0.1, 1, 10]
Gamma_list = [1, 0.1, 0.01]
def try_kernels(kernel_name, c=None, g=None, cm=True):
    print("Fitting model with {} kernal".format(kernel_name))
    if c is not None:
        model = SVC(kernel=kernel_name, C=c, gamma=g)
    else:
        model = SVC(kernel=kernel_name)
    model.fit(X_train, y_train)
    y_pred = model.predict(X_test)
    training_acc=model.score(X_train,y_train)
    validation_acc=model.score(X_test,y_test)
    print("training accuracy with hyperparams:", model.score(X_train,y_train), "\n")
    print("validation accuracy with hyperparams:", model.score(X_test,y_test), "\n")
    if cm:
        cm=metrics.confusion_matrix(y_true=y_test, y_pred=y_pred)
        matplotlib.subplots(figsize=(10, 6))
        sns.heatmap(cm, annot = True, fmt = 'g')
        matplotlib.xlabel("Predicted")
        matplotlib.ylabel("Actual")
        matplotlib.title("Confusion Matrix")
        matplotlib.show()

    print("-----")
    return training_acc, validation_acc
```

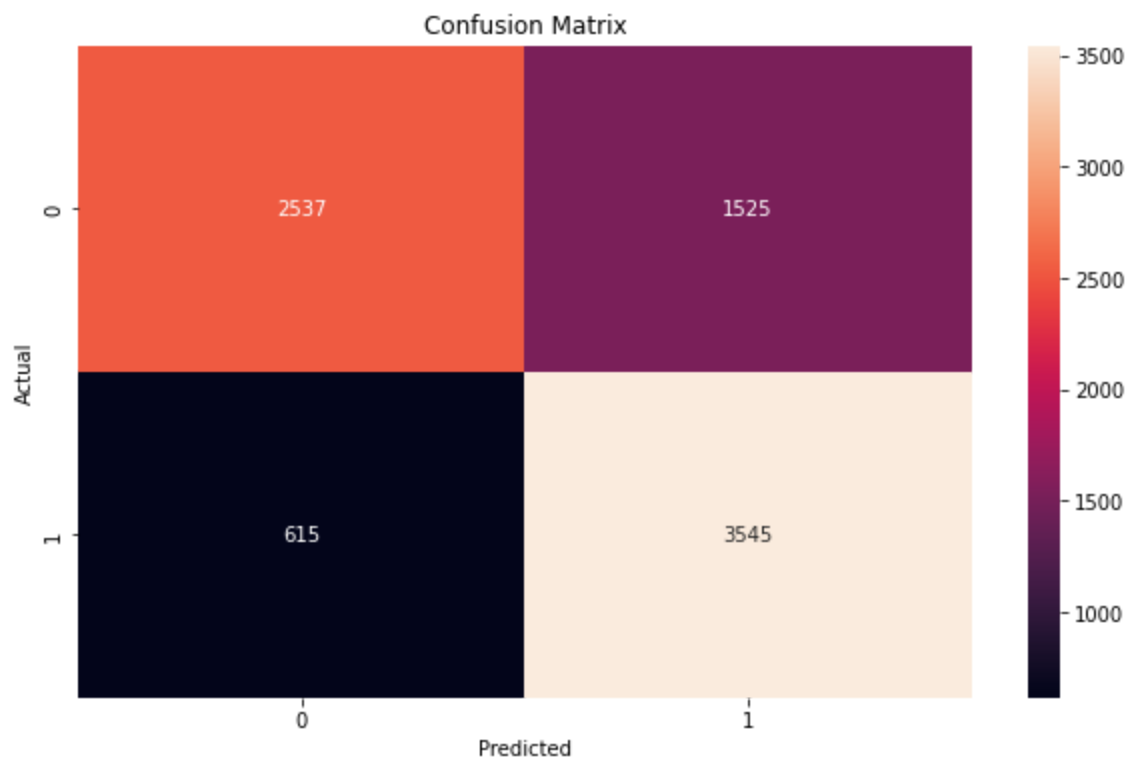
In [24]:

```
kernel_list=["linear", "poly", "sigmoid", "rbf"]
for i in kernel_list:
    try_kernels(i)
```


```
Fitting model with linear kernal
training accuracy with hyperparams: 0.7376637067672221

validation accuracy with hyperparams: 0.7397226952079786
```







Fitting model with poly kernal  
training accuracy with hyperparams: 0.772736487856303  
validation accuracy with hyperparams: 0.7664801751398687



Fitting model with sigmoid kernal  
training accuracy with hyperparams: 0.6351619835380935  
validation accuracy with hyperparams: 0.6365847725614205



Fitting model with rbf kernal  
training accuracy with hyperparams: 0.8092283988160403  
validation accuracy with hyperparams: 0.798710775966918



**Observation:**

The SVM results show that the RBF kernel is performing the best with a training accuracy of 0.809 and validation accuracy of 0.799, followed by the polynomial kernel with a training accuracy of 0.773 and validation accuracy of 0.766. The linear kernel is giving a lower training and validation accuracy of 0.738 and 0.740, respectively. The sigmoid kernel is giving the lowest training and validation accuracy of 0.635 and 0.637, respectively.

The RBF kernel is performing well because it is a non-linear kernel function that can map the input data to a higher-dimensional space, making it easier to separate the classes. It is a popular choice as it can fit complex datasets well. On the other hand, the sigmoid kernel is a simpler kernel function that is not suitable for complex datasets. It is better suited for simpler datasets with linearly separable classes.

In [25]:

```
for c in C_list:
    for g in Gamma_list:
        print("Hyperparamters c and g as",c,g)
        try_kernels(c=c,g=g,cm=False,kernel_name="rbf")
        print()
```

```
Hyperparamters c and g as 0.1 1
Fitting model with rbf kernal
training accuracy with hyperparams: 0.6064955601508333
```

```
validation accuracy with hyperparams: 0.5807589394307954
```

-----

```
Hyperparamters c and g as 0.1 0.1
Fitting model with rbf kernal
training accuracy with hyperparams: 0.7872927056724648
```

```
validation accuracy with hyperparams: 0.787278034541474
```

-----

```
Hyperparamters c and g as 0.1 0.01
Fitting model with rbf kernal
training accuracy with hyperparams: 0.747313789887686
```

```
validation accuracy with hyperparams: 0.7541960593529555
```

-----

```
Hyperparamters c and g as 1 1
Fitting model with rbf kernal
training accuracy with hyperparams: 0.9812269391396018
```

```
validation accuracy with hyperparams: 0.7488445633665775
```

-----

```
Hyperparamters c and g as 1 0.1
Fitting model with rbf kernal
training accuracy with hyperparams: 0.8235413372258038
```

```
validation accuracy with hyperparams: 0.801386523960107
```

-----

```
Hyperparamters c and g as 1 0.01
Fitting model with rbf kernal
training accuracy with hyperparams: 0.772736487856303
```

```
validation accuracy with hyperparams: 0.780345414740939
```

-----

```
Hyperparamters c and g as 10 1
Fitting model with rbf kernal
training accuracy with hyperparams: 0.9997161740258688
```

validation accuracy with hyperparams: 0.7464120651909512

```
-----  
Hyperparamters c and g as 10 0.1  
Fitting model with rbf kernal  
training accuracy with hyperparams: 0.8748327454081012  
  
validation accuracy with hyperparams: 0.7977377766966675  
  
-----
```

```
Hyperparamters c and g as 10 0.01  
Fitting model with rbf kernal  
training accuracy with hyperparams: 0.7925232129100271  
  
validation accuracy with hyperparams: 0.7953052785210412  
  
-----
```

### Observation:

The analysis of the results reveals that the choice of hyperparameters has a significant impact on the performance of the SVM model. For the hyperparameters combination of  $C=0.1$  and  $\gamma=1$ , the model achieved the lowest accuracy, with a training accuracy of 0.6065 and a validation accuracy of 0.5808. In contrast, the highest training accuracy of 0.9997 was achieved with the hyperparameters combination of  $C=10$  and  $\gamma=1$ . However, the corresponding validation accuracy was low, at 0.7464.

I think the best combination of hyperparameters for the RBF kernel appears to be  $c=1$  and  $g=0.1$ , as it achieved the highest validation accuracy of 0.801.

## Random Forest Classifier

In [26]:

```
def perform_random_forest(X_train, y_train, X_val, y_val, X_test, y_test, min_samples_leaf=
    rf_clf = RandomForestClassifier(n_estimators=n_estimators, max_depth=max_depth, min_sar
    rf_clf.fit(X_train, y_train)

    # Evaluate the training and validation accuracy
    train_acc = rf_clf.score(X_train, y_train)
    val_acc = rf_clf.score(X_val, y_val)
    # test_acc = rf_clf.score(X_test, y_test)

    print(f'Training accuracy: {train_acc:.4f}')
    print(f'Validation accuracy: {val_acc:.4f}')
    # print(f'Testing accuracy: {test_acc:.4f}')

    # Analyze feature importance
    if feature_analysis:
        feature_importance = rf_clf.feature_importances_
        sorted_idx = feature_importance.argsort()[::-1]
        print("Feature Importance Ranking:")
        for idx in sorted_idx:
            print(f'Feature {idx+1}: {feature_importance[idx]:.4f}')

    # Generate confusion matrix
    if cm:
        y_pred = rf_clf.predict(X_test)
        cm = confusion_matrix(y_test, y_pred)
        print("Confusion Matrix:")
```

```


matplot.subplots(figsize=(10, 6))
sns.heatmap(cm, annot = True, fmt = 'g')
matplot.xlabel("Predicted")
matplot.ylabel("Actual")
matplot.title("Confusion Matrix")
matplot.show()
return train_acc, val_acc

```

In [27]:

```
perform_random_forest(X_train, y_train, X_val, y_val, X_test, y_test)
```

Training accuracy: 0.9997  
Validation accuracy: 0.8037  
Confusion Matrix:



Out[27]:

```
(0.9997161740258688, 0.8036735190366135)
```

### Observation:

Very High training accuracy indicate that model is overfitting of the data.


In [28]:

```

# Use Random Forest with different values of hyperparameters
perform_random_forest(X_train, y_train, X_val, y_val, X_test, y_test, n_estimators=500, ma

```

Training accuracy: 0.7692  
Validation accuracy: 0.7665  
Confusion Matrix:




Out[28]: (0.7692089364635284, 0.7664517698576816)

### Observation:

The random forest model was trained with hyperparameter tuning by setting the number of estimators to 500 and the maximum depth to 5. The resulting training accuracy was 0.7709 and the validation accuracy was 0.7654. The validation accuracy did not improve compared to the untuned random forest model, but the training accuracy decreased, indicating that the model might be less overfit to the training data.

In [29]: `perform_random_forest(X_train, y_train, X_val, y_val, X_test, y_test, n_estimators=100, ma`

```
Training accuracy: 0.8336
Validation accuracy: 0.7883
Feature Importance Ranking:
Feature 8: 0.2885
Feature 1: 0.1434
Feature 7: 0.1318
Feature 2: 0.0754
Feature 4: 0.0639
Feature 12: 0.0617
Feature 6: 0.0607
Feature 10: 0.0493
Feature 16: 0.0347
Feature 15: 0.0214
Feature 11: 0.0181
Feature 9: 0.0175
Feature 14: 0.0146
Feature 3: 0.0075
Feature 5: 0.0070
Feature 13: 0.0044
Confusion Matrix:
```



Out[29]: (0.833596886023598, 0.788346916433524)

### Observation:

The random forest model with hyperparameter tuning and feature analysis has a training accuracy of 0.8334 and a validation accuracy of 0.7880, which is an improvement over the previous model. The feature importance ranking suggests that feature 8 is the most important, followed by features 7 and 1. Features 4, 2, and 12 also have relatively high importance scores. Features 5, 3, and 13 have the lowest importance scores.

In [30]: `perform_random_forest(X_train, y_train, X_val, y_val, X_test, y_test, n_estimators=50, max`

Training accuracy: 0.7886

Validation accuracy: 0.7772

Out[30]: (0.7886307424076552, 0.7771560637392044)

In [31]: `perform_random_forest(X_train, y_train, X_val, y_val, X_test, y_test, n_estimators=10, max`

Training accuracy: 0.7558

Validation accuracy: 0.7578

Out[31]: (0.7558285691116247, 0.7578153509305438)

In [32]: `perform_random_forest(X_train, y_train, X_val, y_val, X_test, y_test, n_estimators=100, ma`

Training accuracy: 0.8681

Validation accuracy: 0.7960

Out[32]: (0.8681425617321493, 0.7960102177350687)

In [33]: `perform_random_forest(X_train, y_train, X_val, y_val, X_test, y_test, n_estimators=100, ma`

Training accuracy: 0.9997

Validation accuracy: 0.8012

Out[33]: (0.9997161740258688, 0.8012407249726311)

## Observation:

the best performing model is the one with the hyperparameters `n_estimators=100`, `max_depth=100`, `feature_analysis=False`, `max_leaf_nodes=500`, and `min_samples_leaf=5`. This model achieved a training accuracy of 0.8681 and a validation accuracy of 0.7981. While the model with the hyperparameters `n_estimators=100`, `max_depth=None`, `feature_analysis=False`, `max_leaf_nodes=None`, and `min_samples_leaf=1` achieved a higher training accuracy of 0.9997, its validation accuracy of 0.8022 is only slightly better than the best performing model, indicating that it might be overfitting.

## Final individual Models

In [34]:

```
scores=[]
train_acc,val_acc=Logistic_regression(solver="lbfgs",penalty="l1")
lr= LogisticRegression(solver="lbfgs",penalty="l2",max_iter=100,C=1)
scores.append(['Logistic Regression','Saga solver',train_acc,val_acc])
```

Mean training Accuracy: 0.7421643757855898  
Mean validation Accuracy: 0.7404208733730689

In [35]:

```
train_acc,val_acc =try_kernels(c=1,g=0.1,cm=False,kernel_name="rbf")
scores.append(['SVM','rbf kernel',train_acc,val_acc])

rbf_kernel_model = SVC(kernel="rbf",C=1,gamma=0.1, probability=True)
rbf_kernel_model.fit(X_train, y_train)
```

Fitting model with rbf kernal  
training accuracy with hyperparams: 0.8235413372258038  
  
validation accuracy with hyperparams: 0.801386523960107

Out[35]:

```
▼ SVC
SVC(C=1, gamma=0.1, probability=True)
```

In [36]:

```
train_acc,val_acc =perform_random_forest(X_train, y_train, X_val, y_val, X_test, y_test, r
scores.append(['Random Forest','default params',train_acc,val_acc])

rf_clf = RandomForestClassifier(n_estimators=100, max_depth=100,max_leaf_nodes=500,min_sar
rf_clf.fit(X_train, y_train)
```

Training accuracy: 0.8673  
Validation accuracy: 0.7948

Out[36]:

```
▼ RandomForestClassifier
RandomForestClassifier(max_depth=100, max_leaf_nodes=500, min_samples_leaf=5)
```

## Ensemble Classifier

In [37]:

```
#Hard Voting Classifier
def evaluate_accuracy(model):
    model.fit(X_train,y_train)
    t_score = model.score(X_train,y_train)
    print("Accuracy on training data:",t_score)
    p_score = model.score(X_val,y_val)
```



```
print("Accuracy on validation data:", p_score)
return [t_score, p_score]
```

```
model = VotingClassifier(estimators=[('svm', rbf_kernel_model), ('rf', rf_clf), ('lr', lr)], voting='hard')
acc = evaluate_accuracy(model)
scores.append({
    'Voting Classifier':
    'hard',
    acc[0],
    acc[1]
})
```

Accuracy on training data: 0.8322588492884078  
Accuracy on validation data: 0.7936990633742854

In [38]:

```
#Soft Voting Classifier
model = VotingClassifier(estimators=[('svm', rbf_kernel_model), ('rf', rf_clf), ('lr', lr)], voting='soft')
acc = evaluate_accuracy(model)
scores.append({
    'Voting Classifier':
    'soft',
    acc[0],
    acc[1]
})
```

Accuracy on training data: 0.8325832218302721  
Accuracy on validation data: 0.7988079309086485

In [39]:

```
#Stacking
estimator = AdaBoostClassifier(n_estimators=100, learning_rate=0.01)
model = StackingClassifier(estimators=[('svm', rbf_kernel_model), ('rf', rf_clf), ('lr', lr)], voting='hard',
                           base_estimator=estimator)
acc = evaluate_accuracy(model)
scores.append({
    'model': 'Stacking',
    'best params': 'AdaBoost',
    'training accuracy': acc[0],
    'validation accuracy': acc[1]
})
```

Accuracy on training data: 0.8486396626525564  
Accuracy on validation data: 0.8047682763654057

In [40]:

```
#Stacking
estimator = GradientBoostingClassifier(n_estimators=100, learning_rate=0.01, max_depth=100)
model = StackingClassifier(estimators=[('svm', rbf_kernel_model), ('rf', rf_clf), ('lr', lr)], voting='hard',
                           base_estimator=estimator)
acc = evaluate_accuracy(model)
scores.append({
    'model': 'Stacking',
    'best params': 'GradientBoosting',
    'training accuracy': acc[0],
    'validation accuracy': acc[1]
})
```

Accuracy on training data: 0.8588168511535499  
Accuracy on validation data: 0.8044033572558082

In [41]:

```
print(scores)
```

```
[['Logistic Regression', 'Saga solver', 0.7421643757855898, 0.7404208733730689], ['SVM', 'rbf kernel', 0.8235413372258038, 0.801386523960107], ['Random Forest', 'default params', 0.8672910838097555, 0.7947938207030775], {0.8322588492884078, 0.7936990633742854, 'Voting Classifierhard'}, {0.8325832218302721, 0.7988079309086485, 'Voting Classifier', 'soft'}, {'model': 'Stacking', 'best params': 'AdaBoost', 'training accuracy': 0.8486396626525564, 'validation accuracy': 0.8047682763654057}, {'model': 'Stacking', 'best params': 'GradientBoosting', 'training accuracy': 0.8588168511535499, 'validation accuracy': 0.8044033572558082}]
```

```
'validation accuracy': 0.8047682763654057}}, {'model': 'Stacking', 'best params': 'Gradient Boosting', 'training accuracy': 0.8588168511535499, 'validation accuracy': 0.8044033572558082}]
```

### Observation:

I would select gradient boosting classifier as a final model. validation accuracy is almost similar for adaboost and gradient boost but training accuracy is bit higher for gradient boost classifier. Hence going ahead with gradient boost classifier.

## Final Model

In [42]:

```
model.fit(X_train,y_train)
y_pred=model.predict(X_test)

from sklearn.metrics import classification_report

print("model accuracy:",metrics.accuracy_score(y_test, y_pred))
print("model recall:",metrics.recall_score(y_test, y_pred, zero_division=1))
print("model precision:",metrics.precision_score(y_test, y_pred, zero_division=1))
print("classification report:",metrics.classification_report(y_test, y_pred, zero_division=1))
```

```
model accuracy: 0.8068596448552664
model recall: 0.8456730769230769
model precision: 0.7880824372759857
```


```
classification report:                precision    recall  f1-score   support

      0          0.83          0.77          0.80         4062
      1          0.79          0.85          0.82         4160

 accuracy                   0.81         8222
  macro avg              0.81          0.81          0.81         8222
 weighted avg            0.81          0.81          0.81         8222
```

In [44]:

```
cm=metrics.confusion_matrix(y_true=y_test, y_pred=y_pred)
matplot.subplots(figsize=(10, 6))
sns.heatmap(cm, annot = True, fmt = 'g')
matplot.xlabel("Predicted")
matplot.ylabel("Actual")
matplot.title("Confusion Matrix")
matplot.show()
```



## Final Thought

Based on the results, the model has an overall accuracy of 0.8069 and an F1-score of 0.81, which suggests that the model is performing reasonably well. The model also has a recall of 0.8457 and a precision of 0.7881, which indicate that it is able to identify a high proportion of positive instances while also minimizing the false positives.

In terms of future improvements, few suggestions:

1. Having more data may help improve the model's accuracy and generalization performance.
2. Consider exploring additional features or transforming existing features to improve the model's ability to capture patterns in the data.

These two could be possible improvements other than the hyperparameter tuning which we already did in this exercise.